FINAL TEAM STRUCTURE & RESPONSIBILITY MAPPING

---

1. Richa Sinha + Mohit Ranjan — ML Engineers

Primary Ownership (Non-Negotiable)

```
[ AI / ML CORE ]
├── Feature Engineering Logic
│    ├── Offline feature computation
│    ├── Feature validation & contracts
│    └── Feature correctness
│
├── Model Development
│    ├── PyTorch models
│    ├── XGBoost / LightGBM models
│    ├── LLM fine-tuning (HF Transformers)
│    └── Model evaluation & metrics
│
├── Explainability
│    ├── SHAP
│    ├── LIME
│    └── Model confidence scoring
│
└── Model Quality
     ├── Bias detection
     ├── Drift definitions
     └── Performance baselines
```

Tools She Owns

- PyTorch
- Hugging Face Transformers
- XGBoost / LightGBM
- Feast (feature definition layer only)

- SHAP / LIME
- Python ML stack

2. Satyajit + Mohit — Python Engineers

Primary Ownership

```
[ DATA + AI SERVICES ]
├── Data Pipelines
│    ├── Kafka consumers
│    ├── Stream processing logic
│    ├── Batch data jobs
│
├── AI Service Layer
│    ├── FastAPI inference APIs
│    ├── Feature retrieval APIs
│    └── Model invocation wrappers
│
├── Integration Layer
│    ├── Feature Store access
│    ├── Object storage access
│    └── Model registry access
│
└── Async Processing
     ├── Redis Streams
     ├── Retry logic
     └── Backpressure handling
```

Tools He Owns

- Python 3.11
- FastAPI
- Kafka client logic
- Redis Streams
- Feast (online feature access)
- S3 integration
- Model inference glue code

Mohit Ranjan — Systems / DevSecOps / Platform Engineer

This is where the backbone lives.

Primary Ownership

```
[ PLATFORM, INFRA & GOVERNANCE ]
├── Kubernetes Platform
│     ├── Cluster setup
│     ├── Workload isolation
│     ├── Resource scheduling
│     └── Autoscaling
│
├── Streaming Infrastructure
│     ├── Kafka / Redpanda
│     ├── Topic strategy
│     ├── Reliability & replay
│
├── ML Platform
│     ├── Kubeflow
│     ├── Model Registry
│     ├── KServe
│     ├── Triton Inference Server
│
├── Workflow Engines
│     ├── Apache Airflow
│     └── Temporal
│
├── Security & Governance
│     ├── Keycloak (OIDC, RBAC)
│     ├── HashiCorp Vault
│     ├── Policy-as-Code (OPA)
│     └── mTLS & secrets
│
├── Observability
```
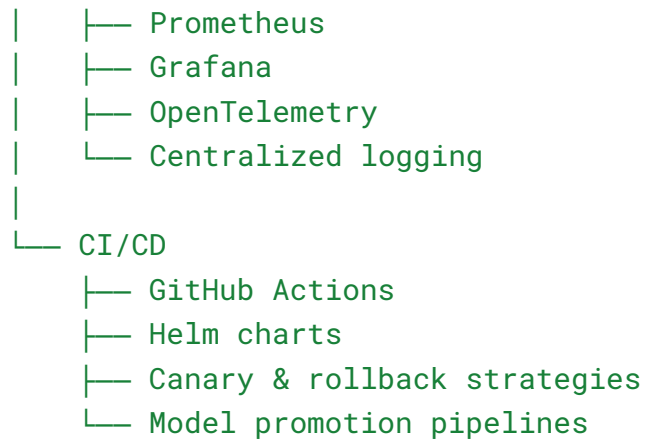
```
|     ├── Prometheus
|     ├── Grafana
|     ├── OpenTelemetry
|     └── Centralized logging
|
└── CI/CD
      ├── GitHub Actions
      ├── Helm charts
      ├── Canary & rollback strategies
      └── Model promotion pipelines
```

[DECISION INTELLIGENCE]
├── Model output interpretation
├── Threshold & confidence calibration
├── Rule–model arbitration
└── Safe-fail logic

Primary: Richa (logic correctness)
Enforcement: Mohit (policy & runtime)